

Parallelized Exact k-Nearest Neighbors Implementation

Rousomanis Georgios (10703)
Department of Electrical and Computer Engineering
Aristotle University of Thessaloniki
Email: rousoman@ece.auth.gr

July 17, 2025

1 Introduction

The goal of this project is to develop an efficient parallel solution to the **approximate all-to-all nearest neighbors (ANN)** problem. This task is crucial in large-scale data analysis applications, where computing exact all-to-all k -nearest neighbors becomes computationally infeasible due to the quadratic complexity in both runtime and memory.

As a first step, we implemented a highly optimized and parallelized **exact k-Nearest Neighbors (k-NN)** algorithm to establish a baseline for correctness and to study the performance characteristics of parallel execution strategies. This implementation allowed us to gain valuable insights into memory management, batching strategies, and parallel workload distribution across multiple threading models.

Building on these foundations, we extended the approach to the more challenging **approximate all-to-all nearest neighbors** problem using a clustering-based approximation strategy. This method reduces the computational cost by partitioning the dataset into smaller subsets and performing local exact k-NN computations, allowing for significant acceleration with a controlled trade-off in accuracy.

Our work systematically evaluates the accuracy, throughput, and scalability of the proposed methods under different parallelization frameworks (Pthreads, OpenMP, OpenCilk), cluster granularities, and thread configurations.

2 k-Nearest Neighbors Implementation

2.1 Problem Definition

The **k-Nearest Neighbors (k-NN)** problem involves finding, for each query point, the k closest points from a reference corpus according to a specified distance metric. Formally, given a set of N data points (the corpus) $C \in \mathbb{R}^{N \times D}$ and a set of M query points $Q \in \mathbb{R}^{M \times D}$, the goal is to identify, for each $q_i \in Q$, the k nearest neighbors in C based on pairwise distances.

In this implementation, the **squared Euclidean distance** is used, which can be efficiently computed in matrix form without explicit loops. The pairwise distance matrix $D \in \mathbb{R}^{M \times N}$ between Q and C is calculated as:

$$D = \|Q\|^2 + \|C\|^2 - 2QC^\top$$

where:

- $\|Q\|^2$ is the vector of squared norms of the query points, broadcasted across columns,
- $\|C\|^2$ is the vector of squared norms of the corpus points, broadcasted across rows,
- QC^\top is the matrix of dot products between query and corpus points.

This formulation avoids explicit nested loops and enables the use of optimized matrix multiplication routines for efficient distance computation.

2.2 Algorithm Overview

2.2.1 Memory Management

Efficient memory management is essential due to the $O(N^2)$ complexity of pairwise distance computations. The algorithm employs the following strategies:

- **Batch Size Estimation:** The batch size B is determined at runtime based on a user-defined fraction of available system memory (e.g., 50%), ensuring intermediate buffers (distance matrices, norms, results) remain within safe memory limits.
- **Batch Processing:** Queries are processed in batches of size B , reducing peak memory usage from $O(N^2)$ to $O(B \times N)$ while enabling efficient blocked matrix operations.
- **Preallocation and Reuse:** Temporary buffers for norms, distances, and neighbor indices are preallocated once and reused across batches, minimizing allocation overhead and avoiding heap fragmentation.

2.2.2 Parallelization Strategy

The main computational bottleneck of the algorithm lies in two stages: (i) the matrix multiplication between the query batches and the corpus, and (ii) the k -selection step to extract the nearest neighbors. To address this, parallelism is applied across both stages as follows:

1. **Workload Distribution:** Each batch of queries is partitioned into contiguous blocks of memory, with each thread responsible for processing a disjoint subset of the queries. This approach ensures cache-friendly memory access and minimizes false sharing.
2. **Dot Product Computation:** Within each batch, each thread independently computes dot products between its assigned query vectors and the full corpus. Matrix multiplication is performed via OpenBLAS in single-threaded mode to avoid thread contention with the outer parallel loop.
3. **k -Selection Step:** After computing pairwise distances, each thread executes a partial sorting procedure (using `QuickSelect`) on its assigned queries to efficiently determine the k nearest neighbors without fully sorting all distances.
4. **Thread Synchronization:** Upon completing the computations for their assigned queries within a batch, all threads synchronize implicitly at the batch boundary. The program then proceeds to the next batch, maintaining strict batch-wise independence to prevent data hazards.

2.2.3 Parallelization Methods

Pthreads Implementation

- A fixed-size thread pool is created explicitly at the start of the program and remains active throughout execution, eliminating the overhead of repeated thread creation and destruction.
- Synchronization is handled manually using mutexes and condition variables to coordinate thread activity.
- Upon program completion, the thread pool along with all auxiliary data structures are properly destroyed to ensure clean resource deallocation.

OpenMP Implementation

- Thread management is fully automated: threads are spawned and managed by the OpenMP runtime, with no need for manual thread pool creation.
- Synchronization and workload distribution are handled implicitly via compiler directives, such as `#pragma omp parallel for`, reducing developer effort and minimizing synchronization code.
- Upon completion of parallel regions, OpenMP automatically manages thread teardown and resource cleanup, requiring no manual intervention for deallocation.

OpenCilk Implementation

- The OpenCilk runtime handles task parallelism dynamically through a work-stealing scheduler, eliminating the need for manual thread management or explicit thread pools.
- Synchronization is implicit, with the runtime ensuring correct execution order via the fork-join model, removing the need for explicit locks or condition variables.
- Task and resource cleanup are fully managed by the OpenCilk runtime, allowing developers to focus solely on expressing parallelism via `cilk_for` without manual resource deallocation.

Aspect	Pthreads	OpenMP	OpenCilk
Thread Management	Manual fixed thread pool	Automatic runtime threads	Dynamic tasks with work-stealing
Synchronization	Manual (mutexes, cond vars)	Implicit via pragmas	Implicit via fork-join
Resource Cleanup	Manual destruction	Automatic cleanup	Automatic cleanup

Table 1: Comparison of Pthreads, OpenMP, and OpenCilk implementations

2.3 Validation of the Algorithm

The validation process ensures the correctness of the implemented k-nearest neighbors algorithm by systematically comparing its output against a reliable reference.

1. **Test Data Generation:** A diverse set of test cases is generated, including both fixed and randomized datasets. For each test case, the ground truth nearest neighbors and distances are computed using a trusted, well-established k-NN implementation. These reference results serve as the standard for correctness.
2. **Algorithm Execution:** The algorithm under test processes each generated query set against the corresponding corpus, producing estimated nearest neighbors and distance values.
3. **Result Comparison:** The output distances and neighbor indices from the tested algorithm are compared element-wise against the reference results within a specified numerical tolerance. Both the distances and indices must match to validate correctness.
4. **Reporting:** Each test case outcome is reported individually as pass or fail. The overall success of the validation is determined by aggregating these results, ensuring comprehensive verification across varying dataset sizes and dimensions.

This procedure rigorously validates the implementation by leveraging precomputed reference results and enforcing strict numerical and index consistency checks.

2.4 Performance of the Algorithm

2.4.1 Benchmarking Procedure

The benchmarking procedure assesses the performance scalability of the k-nearest neighbors implementation using the **MNIST** dataset. The experiments are designed to evaluate how execution time and throughput vary with different parallelization methods and thread counts.

1. The corpus and query matrices are loaded from the MNIST dataset.
2. The algorithm executes the k-NN search for six thread configurations: 1, 2, 4, 8, 16, and 32 threads.
3. Each configuration is evaluated under three parallelization models: Pthreads, OpenMP, and OpenCilk.
4. For every run, the execution time is measured, and throughput is calculated in queries processed per second.

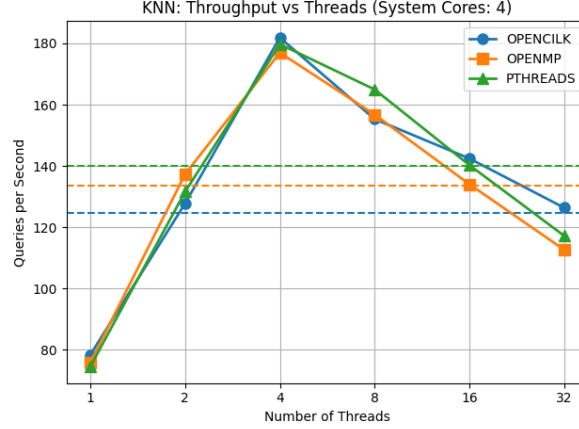


Figure 1: Throughput of the exact k-Nearest Neighbors algorithm as a function of the number of threads used. Solid lines show application-level parallelization with OpenBLAS restricted to a single thread, while dashed lines represent single-threaded application execution with multi-threaded OpenBLAS.

5. All experiments were conducted on a system running Ubuntu LTS 22.04 with a 4-core processor, allowing evaluation of performance scaling both within and beyond the physical core count.

This procedure enables direct comparison of parallelization strategies in terms of efficiency and scalability across varying thread counts.

2.4.2 Benchmark Results

The benchmark results are summarized in Figure 1, which presents the throughput (queries processed per second) of the k-NN implementation as a function of the number of threads used by the application. Separate plots are provided for each of the three parallelization methods: Pthreads, OpenMP, and OpenCilk.

For each method, two curves are depicted:

- A **solid line** shows the throughput when the application itself is parallelized across multiple threads while OpenBLAS is restricted to a single thread. This represents the scenario where all parallel resources are allocated to the k-NN search procedure.
- A **dashed horizontal line** indicates the throughput when the application runs on a single thread while OpenBLAS utilizes all available system cores. This serves as a baseline for the performance attainable through multi-threaded BLAS operations alone, without application-level parallelization.

Across all three parallelization strategies, the solid lines exhibit similar behavior. As the number of application threads increases, the throughput rises, reaching a peak when the thread count matches the number of physical CPU cores (four in our test system). Beyond this point, the throughput begins to decline. This drop is attributed to the overhead introduced by excessive threading, such as increased context-switching, resource contention, and scheduling inefficiencies on the limited number of cores.

Importantly, the peak throughput achieved using the application’s own parallelization (solid line) consistently surpasses the performance achieved by relying solely on OpenBLAS parallelism (dashed line). This demonstrates the effectiveness of explicit workload distribution at the application level, allowing better utilization of compute resources for the specific access patterns and memory requirements of the k-NN search.

Overall, the results highlight both the scalability limits imposed by hardware constraints and the advantages of integrating parallelism directly into the k-NN algorithm, beyond relying solely on low-level numerical libraries.

3 Approximate All-to-All Nearest Neighbors Implementation

3.1 Problem Definition

The **approximate all-to-all nearest neighbors (ANN)** problem involves computing, for each data point in a corpus $C \in \mathbb{R}^{N \times D}$, its k nearest neighbors among all other points in C . This differs from traditional k -NN queries in that the queries and the corpus are identical, and the goal is to approximate the nearest neighbors efficiently.

Due to the $O(N^2)$ computational and memory complexity of exact all-to-all k -NN search, the objective is to develop an approximate solution that reduces runtime and memory footprint, while maintaining acceptable recall accuracy.

In this implementation, we adopt a clustering-based approach to divide the corpus into smaller subsets (clusters), solve exact k -NN within each cluster, and optimize workload distribution through multi-threading.

3.2 Algorithm Overview

The algorithm follows a two-stage pipeline:

- **Stage 1: Clustering.** The data points are partitioned into clusters using a lightweight k -means variant. To ensure cluster quality, small clusters are merged until all clusters contain more than k points, enabling valid intra-cluster k -NN computation.
- **Stage 2: Intra-cluster k -NN.** For each cluster, exact k -NN search is performed internally. Since points are only compared within their assigned clusters, the approach approximates the full k -NN graph.

This pipeline reduces the overall number of distance computations while leveraging parallelization for intra-cluster k -NN computations.

3.2.1 Clustering Strategy

The clustering stage uses a randomized initialization followed by a single assignment step:

- **Initialization:** A fixed number of K_c points are randomly selected from the corpus to serve as initial centroids. The remaining points are assigned to the nearest centroid using exact distance computations.
- **Merging Strategy:** After the initial assignment, clusters with fewer than k points are identified and iteratively merged into their nearest neighboring valid cluster based on centroid distances. This ensures that each cluster has enough points to perform k -NN search locally.
- **Centroid Update:** Following merging, centroids are recomputed by averaging the points within each cluster.

This clustering procedure is executed once prior to k -NN computation, providing a balanced cluster distribution and preventing invalid or degenerate clusters.

3.2.2 Workload Distribution

After clustering, intra-cluster k -NN searches are distributed across threads using a **greedy bin-packing algorithm** based on cluster sizes:

- **Cluster Assignment:** Clusters are sorted by size in descending order. Each cluster is assigned to the thread with the smallest cumulative workload (sum of points across its assigned clusters).
- **Thread Tasks:** Each thread processes its assigned clusters independently, performing exact k -NN within each cluster.
- **Memory Optimization:** Threads dynamically adjust their memory usage based on the number of points they process relative to the total number of points, ensuring adherence to a global maximum memory usage ratio.

This strategy minimizes load imbalance by ensuring that larger clusters are evenly distributed across threads, improving parallel efficiency.

3.2.3 Parallelization Strategy (pros & cons)

The approximate ANN algorithm simplifies parallelization by assigning each thread a fixed subset of clusters only once before computation starts. This static workload distribution offers advantages but also introduces challenges related to load balancing.

- Static assignment avoids repeated thread pool creation and dynamic batching, making the Pthreads implementation simpler—threads are created once, process their assigned clusters independently, then joined at the end.
- However, cluster size variability can cause load imbalance, as some threads may handle more data than others; this contrasts with exact k-NN’s dynamic batching which enables more balanced workload distribution.
- To address imbalance, heuristics like greedy bin-packing are used for cluster assignment, while OpenMP’s static scheduling and OpenCilk’s dynamic work-stealing help redistribute work at run-time.
- Overall, this approach reduces thread management overhead but requires careful workload distribution to maintain parallel efficiency.

3.3 ANN Benchmarks

3.3.1 Benchmarking Procedure

This benchmark assesses the trade-off between accuracy and execution time of the approximate all-to-all nearest neighbors (ANN) algorithm under different configurations, including the number of clusters, thread counts, and parallelization methods. The key steps are:

- Loading and concatenating the training and test subsets of the MNIST dataset.
- Computing the exact all-to-all nearest neighbors on the combined dataset using Euclidean distance as a metric of proximity.
- Executing the ANN algorithm with various cluster counts and parallelization modes, recording throughput and recall by comparison to the exact results.

3.3.2 Recall vs Throughput

The approximate all-to-all nearest neighbors algorithm involves a fundamental tradeoff between accuracy (recall) and computational efficiency (throughput). This tradeoff is primarily governed by the number of clusters used to partition the dataset: fewer clusters yield coarser partitions that tend to retain more true neighbors within the same cluster, improving recall but increasing computation time; conversely, more clusters create finer partitions that reduce workload and increase throughput at the cost of potentially missing some true neighbors.

Figure 2 illustrates how recall degrades as the number of clusters increases, with high accuracy (up to 90%) achieved at low cluster counts (around 5) and recall dropping below 60% as clusters approach 100. This confirms the expected impact of cluster granularity on approximation quality.

The combined effect of this tradeoff on performance is presented in Figure 3, which plots recall against throughput for three parallelization methods—Pthreads, OpenMP, and OpenCilk—across multiple cluster counts ($K_c = 5, 10, 20, 50, 100$). All implementations yield consistent recall values for each cluster count and demonstrate comparable throughput, highlighting the stability of the algorithm across parallel frameworks. This visualization enables users to select an appropriate cluster count to balance speed and accuracy according to their application requirements.

3.3.3 Throughput Scaling with Number of Threads

Figure 4 illustrates throughput scaling with respect to the number of threads for different cluster counts K_c . Across all cases, throughput increases with thread count, peaking around the number of physical CPU cores.

For smaller K_c (e.g., 5, 10, 20), throughput flattens beyond the core count. This occurs because the limited number of large clusters results in significant workload imbalance—execution time is dominated by the largest cluster, leaving extra threads underutilized.

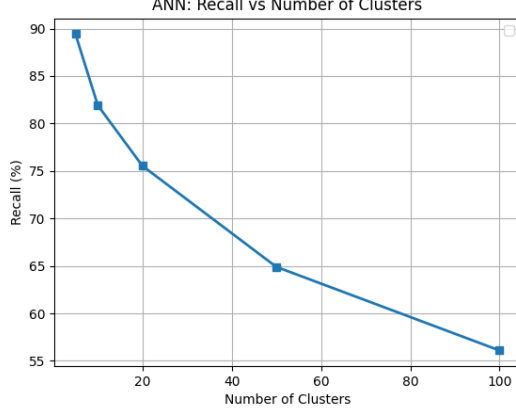


Figure 2: Recall of the approximate all-to-all nearest neighbors algorithm as a function of the number of clusters




Figure 3: Recall versus throughput for the approximate all-to-all nearest neighbors algorithm using Pthreads, OpenMP, and OpenCilk parallelization methods. Points correspond to cluster counts $K_c = 5, 10, 20, 50$, and 100 .

In contrast, for higher K_c (e.g., 50, 100), throughput exhibits a clearer “elbow” behavior: it increases up to the physical core count and then drops slightly. The larger number of clusters improves workload distribution across threads, but the overhead of thread oversubscription (e.g., context switching) prevents further scaling.

Similar trends are observed for OpenMP and OpenCilk, with OpenCilk often providing smoother scaling due to dynamic work stealing.

4 Conclusion and Future Work

In this project, we developed and evaluated parallel algorithms for both exact k-Nearest Neighbors (k-NN) and approximate all-to-all nearest neighbors (ANN) search. Starting with the exact k-NN implementation, we established a high-performance baseline, analyzing its scalability and performance under different parallelization strategies. Building upon this, we introduced a clustering-based approximation scheme for the ANN problem, significantly reducing computational costs while offering tunable trade-offs between accuracy and throughput.

Our benchmarking results demonstrated that:

- The exact k-NN algorithm achieves good scalability up to the number of physical cores, with explicit parallelization outperforming library-only threading.




Figure 4: Throughput vs number of threads for different cluster counts using Pthreads. Throughput increases up to the physical core count and remains stable beyond. Higher cluster counts yield better scaling due to improved workload balance.

- The ANN algorithm offers substantial speedups with controllable recall, where increasing the number of clusters improves throughput at the cost of reduced accuracy.
- Multithreading achieves near-linear speedup up to core count, with OpenCilk providing smoother scaling in scenarios with greater workload imbalance.

Future work will focus on further improving the quality and efficiency of the ANN solution. Key directions include:

- **Merging neighbor lists across clusters** to reduce boundary effects and recover part of the lost accuracy.
- **Exploring advanced clustering techniques**, such as balanced k-means or hierarchical clustering, to achieve more uniform cluster sizes and minimize load imbalance.
- **Investigating hybrid parallelism** (e.g., combining OpenMP with SIMD vectorization) to push performance beyond current thread-level parallelism.
- **Optimizing memory access patterns** for NUMA systems or GPUs to further accelerate computation on modern hardware platforms.

Overall, our work highlights the practicality of clustering-based ANN methods in large-scale nearest neighbor computations and provides a flexible, parallelizable foundation for further extensions.